



Homework 2: Deep Learning

Fabio Cermelli
fabio.cermelli@polito.it

Originally made by
Antonio D'Innocente




Overview

- You are going to train a Convolutional Neural Network for image classification
 - AlexNet on Caltech-101
- Start from a provided template code, and implement experiments

AlexNet ([link](#))

5 convolutional layers
3 fully connected layers

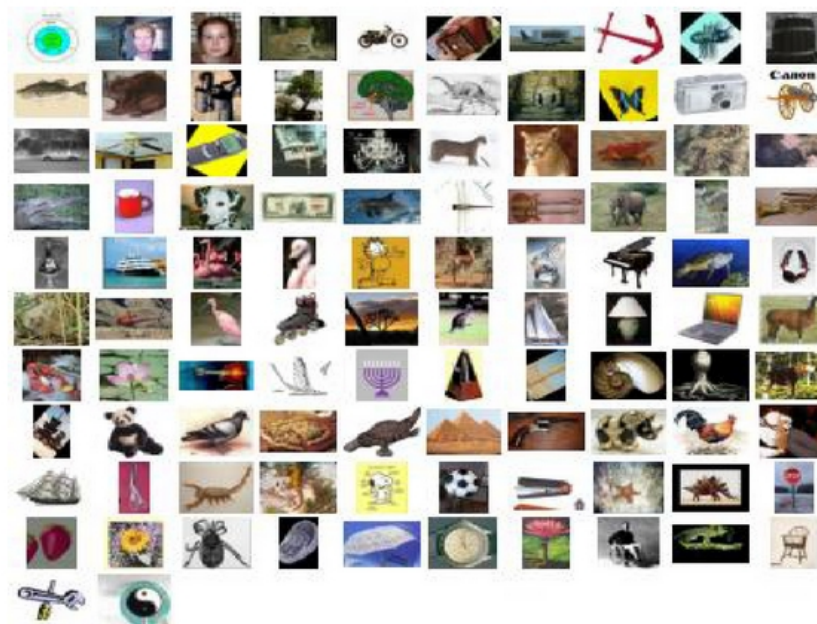


Caltech-101 ([link](#))

101 Classes (+ background)

Common Categories

Used dataset



GPUs

Deep Learning needs GPUs
acceleration

We use Google Colab

- <https://colab.research.google.com/>





Code Template

- Starting code is available at:
 - <https://colab.research.google.com/drive/1495rzqiMxfqwqBjaanDwDGN7eOS4oXRX>
 - Save a copy on your own drive “File -> Save a Copy on Drive” and start working on it
 - “Edit -> Notebook Settings”: Check that hardware acceleration is using GPU
 - If you find it more convenient, work and debug locally with CPU (remember to set the argument `DEVICE='cpu'`), and upload working code when ready for training
- Dataset available via github repository:
 - <https://github.com/MachineLearning2020/Homework2-Caltech101>



Homework 2

- 0- First Step
- 1- Data Preparation
- 2- Training from scratch
- 3- Transfer Learning
- 4- Data Augmentation
- 5- (Extra) Beyond AlexNet



0- First Step

A. Study code and data

- Read the comments, understand what everything is doing
- Pay extra attention to: Preprocessing, Datasets, Training, Validation, Testing
- Also explore the data provided on github, and study the ImageFolder class of `torchvision.datasets`

B. Run the code

- Training should take less than ten minutes
- You have to stay connected



1- Data Preparation

- Create the dataset for Caltech, following the code provided in the GitHub repository and taking as reference the ImageFolder dataset.
- train/test splits are actually provided in the github repository
- There is also a BACKGROUND folder in Caltech-101 that is used but shouldn't be



1- Data Preparation

- A. Use the provided train/test splits for training dataset and test dataset
- B. Filter out the BACKGROUND class (which is incorrectly provided also in the splits)

There is not an unique way, however the most elegant solution is an ad-hoc dataset class which reads the split files. A template is provided in the github repository with the dataset



2- Training from scratch

- For training, we need a validation set for hyperparameter tuning and model selection
- In deep learning, it is often prohibitive to do k-fold cross validation, due to time required to complete trainings
- The most common procedure is to have a single validation set for both hyperparameter tuning and model selection

video: training and validation must be of the same size



2- Training from scratch

- A. Split the training set in training and validation sets
 - The validation set must be half the size of the training set
 - Be careful to not filter out entire classes from the sets!
 - Aim for half samples of each class in train and the other half in val
- B. Implement model selection with validation
 - After each training epoch, evaluate (= test) the model on the validation set
 - Use only the best performing model on the validation set for testing



2- Training from scratch

The current implementation trains using SGD with momentum for 30 epochs with an initial learning rate (LR) of 0.001 and a decaying policy (STEP_SIZE) after 20 epochs after 20 epochs the learning rate is decreased by a value (ex. 1/10 of initial LR)

- C. Experiment with at least two different sets of hyperparameters (suggestion: experiment changing LR and one among: decaying policy, optimiser, epochs)
- The first hyperparameter to optimize is the learning rate
 - Low learning rate = the model converges too slowly (loss decreases slowly)
 - High learning rate = the model converges fast but accuracy is sub-optimal
 - Too high learning rate = the model diverges (loss increases)



3- Transfer Learning

- Deep Learning needs very large datasets to train good features
- Caltech-101 = very small dataset
- Solution: Transfer Learning (finetuning)
 - Use the weights learned by training on a large related dataset as a starting point for training on the small dataset



3- Transfer Learning

- A. Load AlexNet with weights trained on the ImageNet dataset
 - (read the documentation)
- B. Change the Normalize function of Data Preprocessing to Normalize using ImageNet's mean and standard deviation
 - (also search for this)
- C. Run experiments with at least three different sets of hyperparameters

LR, weight decay, epochs ...



3- Transfer Learning

- When transfer learning, we can also “freeze” part of the network and only train certain layers (not always the best choice, but faster and can prevent overfitting in some cases)

With the best hyperparameters you obtained in the previous step:

- D. Experiment by training only the fully connected layers (freeze other layers) => freeze all convolutional
- E. Experiment by training only the convolutional layers => freeze fully connected
 - Compare all results you obtained



4- Data Augmentation

- Deep Learning leverages huge amounts of data
- It is common to artificially increase the dataset size by applying label-preserving transformations to images
 - The image changes, but the label is the same
 - Prevents overfitting by increasing cardinality of data
- Examples of data augmentation include:
 - Cropping, Horizontal Flipping
- The current preprocessing for our input data is an Image Resize followed by a 224x244 central crop



4- Data Augmentation

- A. Apply at least three different sets of preprocessing for training images
- <https://pytorch.org/docs/stable/torchvision/transforms.html>
 - no need to use complex transformations, such as RandomAffine
 - if test data is very similar to training data, some data augmentation policies may worsen accuracy
 - some transforms (FiveCrop and TenCrop) are designed for evaluation only
 - if the loss keeps decreasing, increase training epochs or increase learning rate

Test-time data augmentation is less common in learning and research (but is used in competition). If you feel like it, try to integrate the TenCrop transform for test data (not trivial)



5- (Extra) Beyond AlexNet

- torchvision implements a lot of convolutional neural networks
 - <https://pytorch.org/docs/stable/torchvision/models.html>
- You can experiment with different models (like ResNets or VGGs)
 - some ConvNets need an input size higher than 224x224
 - Larger models consume GPU memory, you might need to decrease the batch size (a lot)
 - Larger models need much longer to train and Colab's K80s are not the fastest GPU around (freezing layers is an option)
 - Larger models are more accurate



Submission rules

- Deadline: **Two week before the first round (NOT the first you enrol in).**
- Uploading: through “Portale della didattica”

Submit a zip named <YOUR_ID>_homework2.zip. The zip should contain two items:

- A pdf report describing data, your implementation choices, results and discussions
- Code



FAQ

1. *Should I submit the .ipynb file or the .py?*
 - Submit both
2. *Should the size of the training and validation sets be equal, or should they be $\frac{2}{3}$ - $\frac{1}{3}$?*
 - They were meant to be equal (50-50), however it is fine to go for $\frac{2}{3}$ - $\frac{1}{3}$ if the point raised ambiguity
 - The important thing is to respect class balancing
3. *Should points 4 and 5 be implemented with transfer learning?*
 - Yes
4. *Should a criterion be used for searching hyperparameters?*
 - It is not required but it is preferable to adopt some kind of criterion and motivate it ex. gridsearch
5. *Should multiple code notebooks be provided (e.g. one for each point)?*
 - Provide just the last implementation, leave previous points as comments



FAQ

6. *Which hyperparameters can we optimize?*

- Learning rate, number of epochs and step down policy are good candidates
- Batch size and weight decay are also options

7. *Should results be fully deterministic?*

- No. If you observe significant deviation between experiments, do multiple runs and report average and std

8. *I have used hyperparameters X/Y and the result is Z, is this ok?*

- Annotate the results you obtained and write them in the final report

9. *My answer for point A is this, is it ok?*

- Provide your implementation in the code, you will get feedback after the hard deadline



FAQ

10. Technical issues (colab notebook, python errors)

- Many problems with colab can be solved by restarting the runtime (Runtime -> Restart and run all)
- Python bugs, pytorch errors, colab issues and other problems are considered as part of the assignment